

▼ 1. Download the dataset

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import skew
%matplotlib inline

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import Ridge
from sklearn.svm import SVR
import warnings
```

▼ 2. Load the dataset into the tool.

```
data = pd.read_csv('/content/abalone.csv')
```

```
data['age'] = data['Rings'] + 1.5
data.drop('Rings', axis = 1, inplace = True)
```

```
print('This dataset has {} observations with {} features.'.format(data.shape[0], data.shape[1]))
```

This dataset has 4177 observations with 9 features.

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Sex             4177 non-null   object
 1   Length          4177 non-null   float64
 2   Diameter        4177 non-null   float64
 3   Height          4177 non-null   float64
 4   Whole weight    4177 non-null   float64
 5   Shucked weight  4177 non-null   float64
 6   Viscera weight  4177 non-null   float64
```

```

7  Shell weight      4177 non-null    float64
8  age               4177 non-null    float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB

```

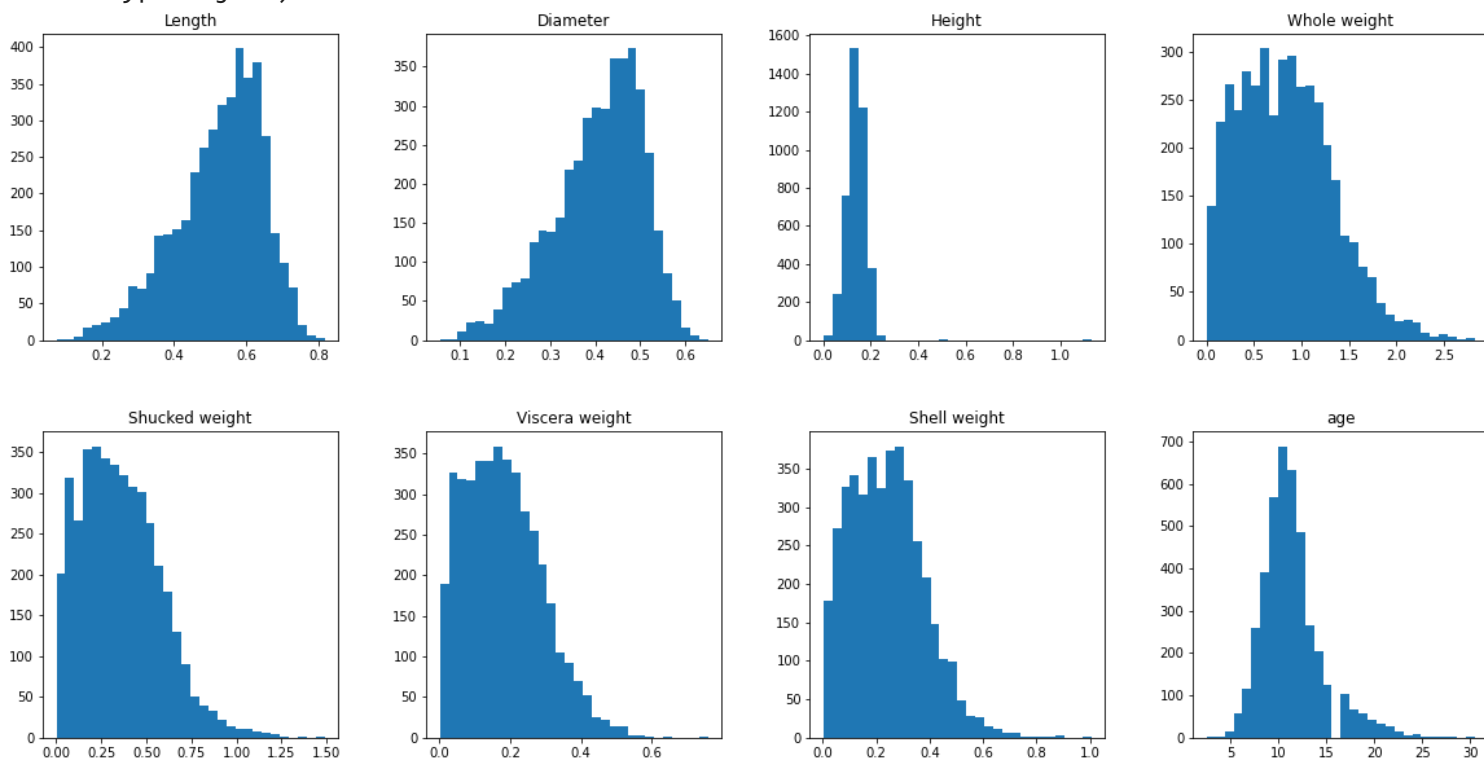
Perform Below Visualizations:Univariate Analysis, Bi-Variate Analysis, Multi-Variate Analysis

```
data.hist(figsize=(20,10), grid=False, layout=(2, 4), bins = 30)
```

```

array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f87728c6f10>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f87723ef1d0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f87724076d0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f87723b5bd0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f8772378110>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f877232e610>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f8772ebe3d0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f8772cfb890>]],
      dtype=object)

```



```
numerical_features = data.select_dtypes(include=[np.number]).columns
categorical_features = data.select_dtypes(include=[np.object]).columns
numerical_features
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `np.object` is deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0
```


```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'age'],
      dtype='object')
```

```
categorical_features
```

```
Index(['Sex'], dtype='object')
```


```
sns.countplot(x = 'Sex', data = data, palette="Set3")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8771e294d0>
```



```
plt.figure(figsize = (20,7))
sns.swarmplot(x = 'Sex', y = 'age', data = data, hue = 'Sex')
sns.violinplot(x = 'Sex', y = 'age', data = data)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 56.2% of the poi
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 52.2% of the poi
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 58.5% of the poi
warnings.warn(msg, UserWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7f8771df52d0>
```

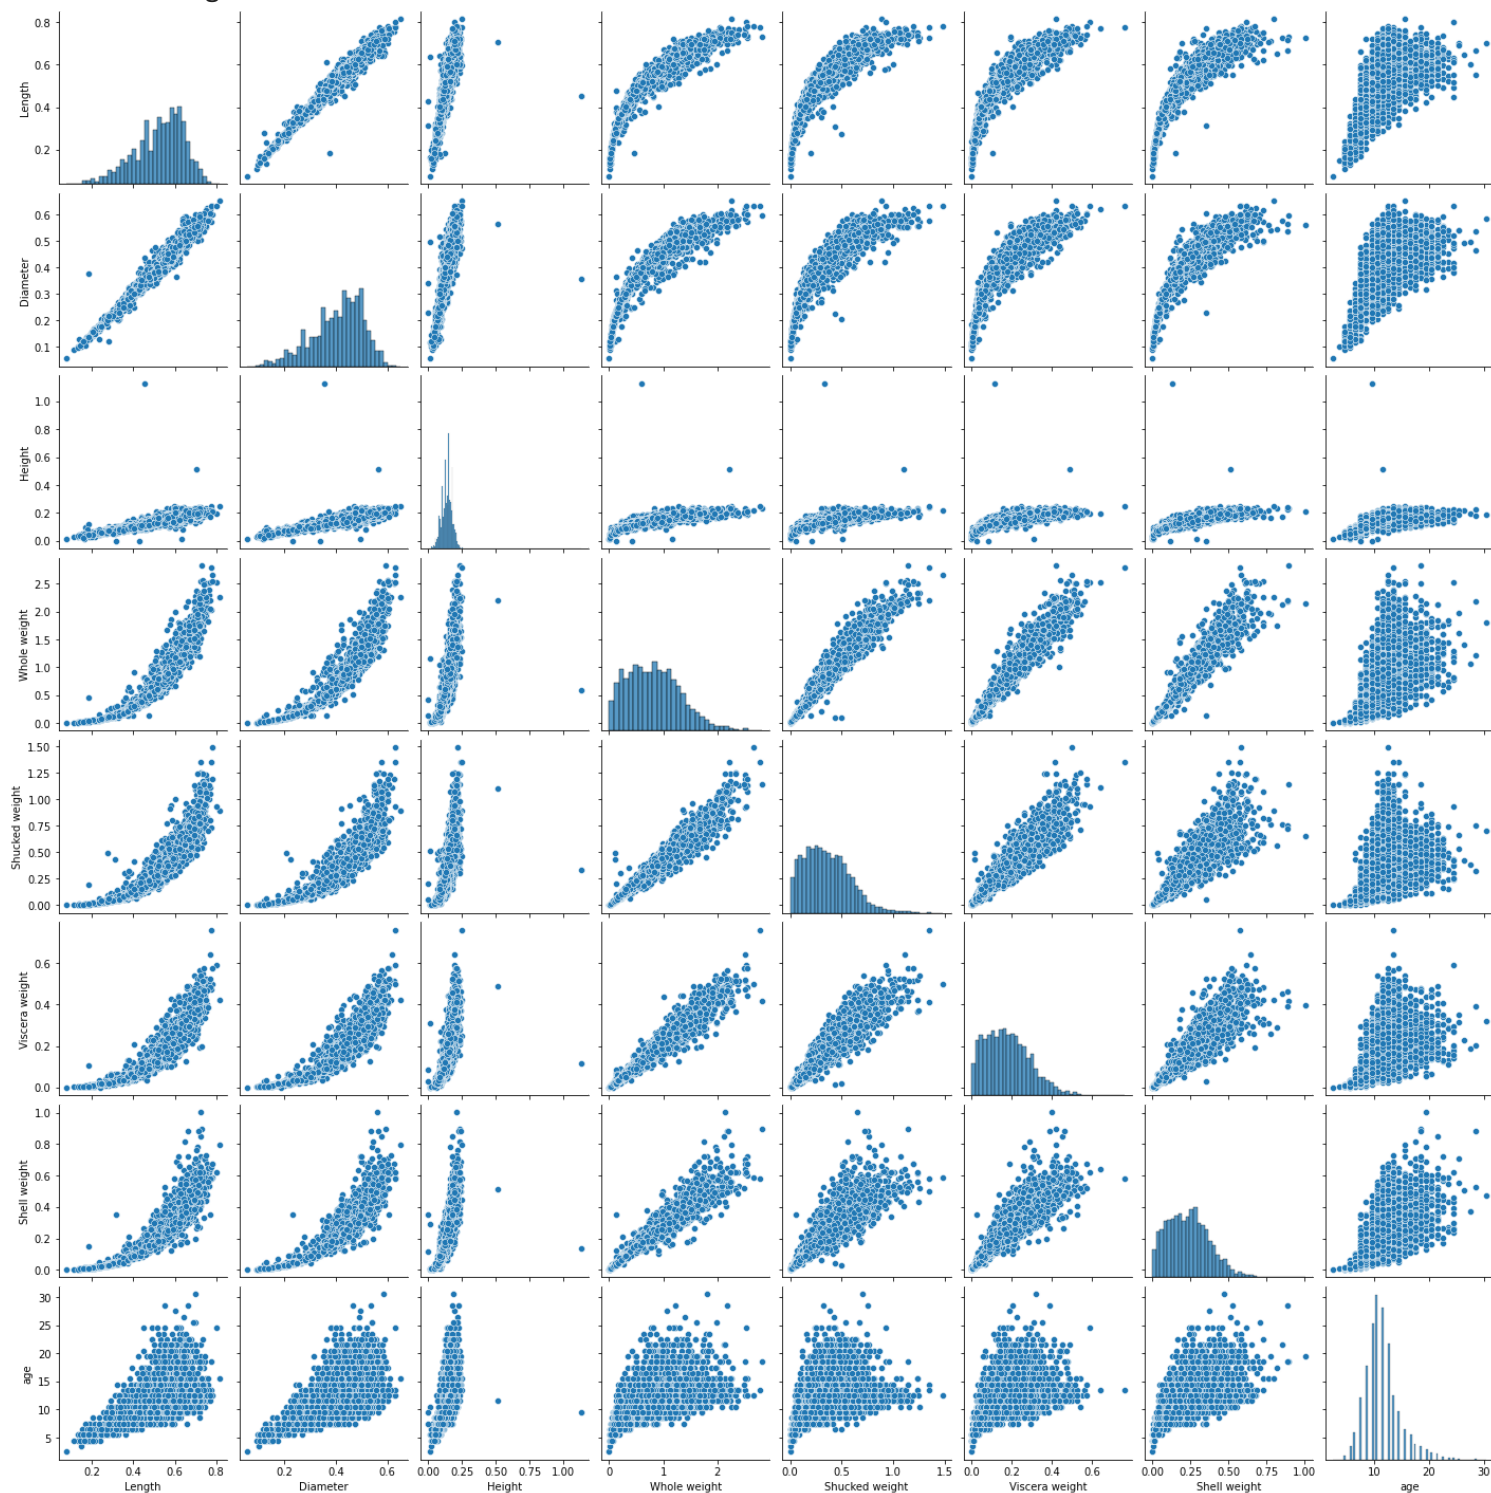


```
data.groupby('Sex')[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
'Viscera weight', 'Shell weight', 'age']].mean().sort_values('age')
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age
Sex								
I	0.427746	0.326494	0.107996	0.431363	0.191035	0.092010	0.128182	9.390462
M	0.561391	0.439287	0.151381	0.991459	0.432946	0.215545	0.281969	12.205497
F	0.579093	0.454732	0.158011	1.046532	0.446188	0.230689	0.302010	12.629304


```
sns.pairplot(data[numerical_features])
```

```
<seaborn.axisgrid.PairGrid at 0x7f2d21671dd0>
```

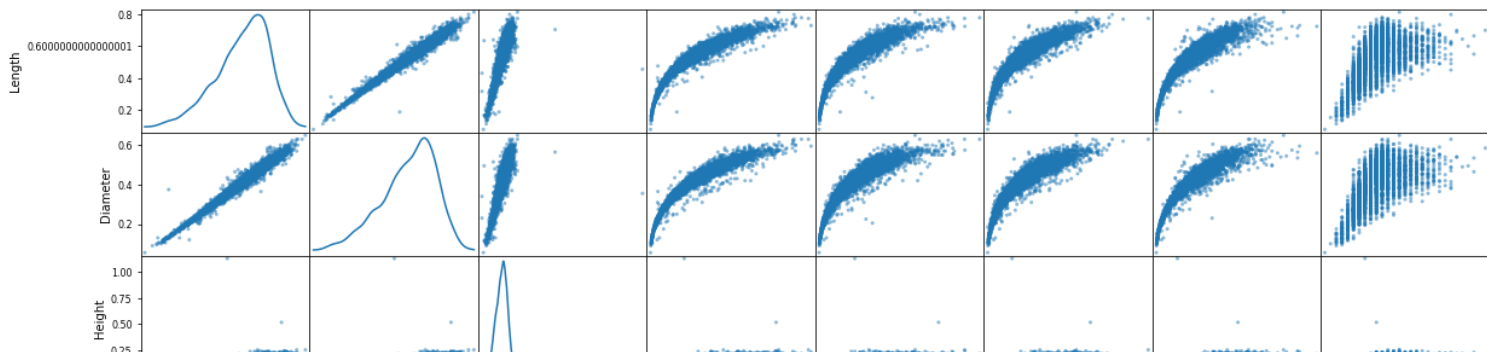


```
plt.figure(figsize=(20,7))  
sns.heatmap(data[numerical_features].corr(), annot=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f876f48eb50>



```
pd.plotting.scatter_matrix(data.loc[:, 'Sex':'age'], diagonal="kde",figsize=(20,15))
plt.show()
```



```
ax = data[['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
          'Viscera weight', 'Shell weight', 'age']].plot(figsize=(20,15))
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5));
```

4. Perform descriptive statistics on the dataset

```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	41
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	

```
data['Sex'].describe()
```

```
count      4177
unique        3
top          M
freq       1528
Name: Sex, dtype: object
```

```
data['Sex'].value_counts()
```

```
M      1528
I      1342
F      1307
Name: Sex, dtype: int64
```

```
data.kurtosis()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance
    """Entry point for launching an IPython kernel.
Length      0.064621
Diameter    -0.045476
Height      76.025509
Whole weight -0.023644
Shucked weight 0.595124
Viscera weight 0.084012
Shell weight 0.531926
age         2.330687
dtype: float64
```



```
skew_values = skew(data[numerical_features], nan_policy = 'omit')
dummy = pd.concat([pd.DataFrame(list(numerical_features), columns=['Features']),
                    pd.DataFrame(list(skew_values), columns=['Skewness degree'])], axis = 1)
dummy.sort_values(by = 'Skewness degree' , ascending = False)
```

	Features	Skewness degree
2	Height	3.127694
7	age	1.113702
4	Shucked weight	0.718840
6	Shell weight	0.620704
5	Viscera weight	0.591640
3	Whole weight	0.530768
1	Diameter	-0.608979
0	Length	-0.639643

5.Check for Missing values and deal with them.

```
missing_values = data.isnull().sum().sort_values(ascending = False)
percentage_missing_values = (missing_values/len(data))*100
pd.concat([missing_values, percentage_missing_values], axis = 1, keys= ['Missing values', '% Missing'])
```

	Missing values	% Missing
Sex	0	0.0
Length	0	0.0
Diameter	0	0.0
Height	0	0.0
Whole weight	0	0.0
Shucked weight	0	0.0
Viscera weight	0	0.0
Shell weight	0	0.0
age	0	0.0

```
data.head
```

	<bound	method	NDFrame.head of			Sex	Length	Diameter	Height	Whole weight	Shucked weight	\
0	M	0.455	0.365	0.095			0.5140		0.2245			
1	M	0.350	0.265	0.090			0.2255		0.0995			
2	F	0.530	0.420	0.135			0.6770		0.2565			

```

3      M    0.440    0.365    0.125    0.5160    0.2155
4      I    0.330    0.255    0.080    0.2050    0.0895
...    ..    ...    ...    ...    ...    ...
4172   F    0.565    0.450    0.165    0.8870    0.3700
4173   M    0.590    0.440    0.135    0.9660    0.4390
4174   M    0.600    0.475    0.205    1.1760    0.5255
4175   F    0.625    0.485    0.150    1.0945    0.5310
4176   M    0.710    0.555    0.195    1.9485    0.9455

```

```

      Viscera weight  Shell weight  age
0          0.1010      0.1500  16.5
1          0.0485      0.0700   8.5
2          0.1415      0.2100  10.5
3          0.1140      0.1550  11.5
4          0.0395      0.0550   8.5
...          ...          ...   ...
4172         0.2390      0.2490  12.5
4173         0.2145      0.2605  11.5
4174         0.2875      0.3080  10.5
4175         0.2610      0.2960  11.5
4176         0.3765      0.4950  13.5

```

```
[4177 rows x 9 columns]>
```


▼ 6.Find the outliers and replace them outliers

```

original_data = data.copy()
data = pd.get_dummies(data)
dummy_data = data.copy()
data.boxplot( rot = 90, figsize=(20,5))


```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f876c4c0590>
```




```
var = 'Viscera weight'
```

```
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
# outliers removal
data.drop(data[(data['Viscera weight']> 0.5) & (data['age'] < 20)].index, inplace=True)
data.drop(data[(data['Viscera weight']<0.5) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Shell weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
data.drop(data[(data['Shell weight']> 0.6) & (data['age'] < 25)].index, inplace=True)
data.drop(data[(data['Shell weight']<0.8) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Shucked weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
data.drop(data[(data['Shucked weight']>= 1) & (data['age'] < 20)].index, inplace=True)
data.drop(data[(data['Shucked weight']<1) & (data['age'] > 20)].index, inplace=True)
```

```
var = 'Whole weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Whole weight']>= 2.5) & (data['age'] < 25)].index, inplace=True)
data.drop(data[(data['Whole weight']<2.5) & (data['age'] > 25)].index, inplace=True)
```


```
var = 'Diameter'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Diameter']<0.1) & (data['age'] < 5)].index, inplace=True)
data.drop(data[(data['Diameter']<0.6) & (data['age'] > 25)].index, inplace=True)
data.drop(data[(data['Diameter']>=0.6) & (data['age']< 25)].index, inplace=True)
```


```
var = 'Height'
```

```
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Height']>0.4) & (data['age'] < 15)].index, inplace=True)
data.drop(data[(data['Height']<0.4) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Length'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Length']<0.1) & (data['age'] < 5)].index, inplace=True)
data.drop(data[(data['Length']<0.8) & (data['age'] > 25)].index, inplace=True)
data.drop(data[(data['Length']>=0.8) & (data['age']< 25)].index, inplace=True)
```

▼ 7. Check for Categorical columns and perform encoding.

```
from sklearn import preprocessing
label = preprocessing.LabelEncoder()
```

```
original_data['Sex']= label.fit_transform(original_data['Sex'])
```

```
original_data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

8. Split the data into dependent and independent variables.

```
X = original_data.drop('age', axis = 1)
y = original_data['age']
```

X

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550
...
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490
4173	2	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605
4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950

4177 rows × 8 columns

y

0	16.5
1	8.5
2	10.5
3	11.5
4	8.5
...	...
4172	12.5

```
4173    11.5
4174    10.5
4175    11.5
4176    13.5
Name: age, Length: 4177, dtype: float64
```

▼ 9. Scale the independent variables

```
# Normalized Y
```

```
from sklearn import preprocessing
Y=y.values.reshape(-1,1)
```

```
normalized_Y = preprocessing.normalize(Y)
```

```
print (normalized_Y)
```

```
[[1.]
 [1.]
 [1.]
 ...
 [1.]
 [1.]
 [1.]]
```

```
# Standardized Y
```

```
standard_Y = Y.copy()
```

```
from sklearn import preprocessing
```

```
ss = preprocessing.StandardScaler()
ss.fit(standard_Y)
```

```
print (standard_Y)
```

```
[[16.5]
 [ 8.5]
 [10.5]
 ...
 [10.5]
 [11.5]
 [13.5]]
```

▼ 10. Split the data into training and testing

```
X = data.drop('age', axis = 1)
y = data['age']
```

```

standardScale = StandardScaler()
standardScale.fit_transform(X)

selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)

X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)

```

▼ 11. Build the Model

12. Train the Model

```

np.random.seed(10)
def rmse_cv(model, X_train, y):
    rmse =- (cross_val_score(model, X_train, y, scoring='neg_mean_squared_error', cv=5))
    return(rmse*100)

models = [LinearRegression(),
          Ridge(),
          SVR(),
          RandomForestRegressor(),
          GradientBoostingRegressor(),
          KNeighborsRegressor(n_neighbors = 4),]

names = ['LR', 'Ridge', 'svm', 'GNB', 'RF', 'GB', 'KNN']

for model,name in zip(models,names):
    score = rmse_cv(model,X_train,y_train)
    print("{}      : {:.6f}, {:.4f}".format(name,score.mean(),score.std()))

    LR      : 348.361383, 25.766755
    Ridge    : 353.619706, 31.391040
    svm      : 374.162503, 40.223080
    GNB      : 343.729821, 22.225674
    RF       : 338.530444, 21.864369
    GB       : 382.908726, 26.213802

def modelfit(alg, dtrain, predictors, performCV=True, printFeatureImportance=True, cv_folds=5):
    #Fit the algorithm on the data
    alg.fit(dtrain[predictors], dtrain['age'])

    #Predict training set:
    dtrain_predictions = alg.predict(dtrain[predictors])
    #dtrain_predprob = alg.predict_proba(dtrain[predictors])[:,1]

    #Perform cross-validation:
    if performCV:
        cv_score = -cross_val_score(alg, dtrain[predictors], dtrain['age'], cv=cv_folds,
                                    scoring='r2')

    #Print model report:

```



```

print ("\nModel Report")
print( "RMSE : %.4g" % mean_squared_error(dtrain['age'].values, dtrain_predictions))
print( "R2 Score (Train): %f" % r2_score(dtrain['age'], dtrain_predictions))

if performCV:
    print( "CV Score : Mean - %.7g | Std - %.7g | Min - %.7g | Max - %.7g" % (np.mean(cv_score),np.
                                                                              np.min(cv_score),np.ma

#Print Feature Importance:
if printFeatureImportance:
    feat_imp = pd.Series(alg.coef_, predictors).sort_values(ascending=False)
    plt.figure(figsize=(20,4))
    feat_imp.plot(kind='bar', title='Feature Importances')
    plt.ylabel('Feature Importance Score')

```

▼ 13. Test the Model


14. Measure the performance using Metrics

```

# Base Model
predictors = [x for x in data.columns if x not in ['age']]
lrm0 = Ridge(random_state=10)
modelfit(lrm0, data, predictors)

```

Model Report
 RMSE : 3.593
 R2 Score (Train): 0.529894
 CV Score : Mean - -0.4503433 | Std - 0.08079434 | Min - -0.514565 | Max - -0.3061263



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 10:20 AM

